

List of elements used:

- Feature engineering of text data
- Four supervised learning models: Ridge Regression, Naive Bayes, Random Forest, and SVM
- Development and tuning of a deep learning model
- Analysis of feature importance
- Methods to address class imbalance (used the CountVectorizer to help with class imbalance, also used class weights in the training of the ANN so that it wouldn't favor the Dinner class so heavily.

Also, the ridge regression has built-in cross validation that I tuned, but I'm not certain if that counts for the implementation of advanced cross-validation.